


KNOWLEDGEHUB-AI SYSTEM ARCHITECTURE DIAGRAM

Figure 1: High-Level Vector Ingestion and RAG Query Pipeline



Key Architectural Specifications:

- Hybrid Search Strategy: Dense vector similarity combined with BM25 lexical search.
- Latency Benchmark: Sub-120ms retrieval delay for top\_k=5 documents.
- Target Diagram ID: FIG-ARCH-2026-V1